



# MONGOOSE 2.0

# WHAT IS MONGOOSE

Mongoose is a benchmarking tool initially designed for cloud storage performance testing

- 1M of concurrent connections
- 1M of operations per second
- 1M of items which may be processed multiple times in the circular load mode
- 1M of items which may be stored in the storage mock

# TOP FEATURES

1. Distributed Mode
2. Rich Metrics Reporting
3. Different operation types  
*(Create, Update, Append, Read, Delete)*
4. Abstract Load Engine  
*(load with objects, files, containers, directories, etc)*
5. Cloud Storages support  
*(S3, Atmos, Swift)*
6. Flexible Load Limitation  
*(by count, time, rate)*
7. Custom Content Generation and Verification
8. Circular Load Mode
9. Dynamic Configuration Parameters
10. Custom Item Naming Schemes

# CURRENT USABILITY ISSUES

- Not enough flexible  
*requires Java programming to implement a custom scenario*
- Error-prone and complicated scenario configuration  
*the usual way to run from the CLI looks like:*

```
java
-Dload.server.addrs=10.248.236.69,10.248.236.68,10.248.236.67,10.248.
236.66
-Dstorage.addrs=10.247.235.65,10.247.235.64,10.247.235.63,10.247.235.
62 -Dload.threads=100 -Ddata.size=16MB -Drun.id=mySimpleReadTest1
-Dapi.type.s3.bucket=bucket1 -Dscenario.type.single.load=read
-Ditem.src.fpath=mongoose-1.4.0/log/mySimpleWriteTest1/items.csv -jar
mongoose-1.4.0/mongoose.jar client
```

# NEW REQUIREMENTS


- Be able to execute custom scenarios
- Make the tool use-case oriented
- Support various mixed load cases
- Support weighted load case
- Include rich set of example scenarios into the distribution

# THE APPROACH

- Use JSON scenario files on input
- Aggregate the load jobs into the job containers
- Make job containers executable in parallel and sequentially
- Make job containers and single load jobs configurable individually

# TRADE-OFFS

- No backward compatibility with 1.x versions
- GUI for 2.x is not ready yet



# HOW TO RUN A SCENARIO FILE

- Specify the scenario file as a CLI argument:

```
java [<DEFAULTS_OVERRIDING>] -jar mongoose.jar [<MODE>] -f  
<PATH/TO/SCENARIO.json>
```

- Or pass the scenario content on the standard input:

```
cat <PATH/TO/SCENARIO.json> | java [<DEFAULTS_OVERRIDING>]  
-jar mongoose.jar [<MODE>]
```



# THE SCENARIO FILE OVERVIEW

The root node is always a job container:

```
{  
    "type" : <JOB_CONTAINER_TYPE>  
    . . .  
    . . .  
}
```

# JOB CONTAINER TYPES

Load	<ul style="list-style-type: none"><li>• Single load job</li><li>• Cannot include another job containers</li></ul>
Precondition	Same as "Load" but doesn't produce the metrics output CSV files
Rampup	<ul style="list-style-type: none"><li>• Multiple load jobs</li><li>• Cannot include another job containers</li></ul>
Parallel	<ul style="list-style-type: none"><li>• Should include other job containers</li><li>• Nested job containers are executed <i>in parallel</i></li></ul>
Sequential	<ul style="list-style-type: none"><li>• Should include other job containers</li><li>• Nested job containers are executed <i>sequentially</i></li></ul>
Command	Execute a shell command

# JOB CONTAINER CONFIGURATION

- Any job container can contain an optional "config" node:

```
{  
  "type" : <JOB_CONTAINER_TYPE>  
  "config" : {  
    // here are the configuration hierarchy  
  }  
}
```

- The layout of the "config" subtree is the same as for default configuration

# CONFIGURATION EXAMPLE

```
{  
  "type" : "load"  
  "config" : {  
    "storage" : {  
      "addrs" : [  
        "192.168.0.1", "192.168.0.2", "192.168.0.3"  
      ]  
    }  
  }  
}
```

# SEQUENTIAL JOBS EXECUTION EXAMPLE

```
{  
  "type" : "sequential"  
  "jobs" : [  
    {  
      // 1st job container  
    }, {  
      // 2nd job container  
    }  
  ]  
}
```

# PARALLEL JOBS EXECUTION EXAMPLE

```
{  
  "type" : "parallel"  
  "jobs" : [  
    {  
      // 1st job container  
    }, {  
      // 2nd job container  
    }  
  ]  
}
```

# SHELL COMMAND JOB EXAMPLE

```
...  
"jobs" : [  
  {  
    ...  
  }, {  
    "type" : "command",  
    "value" : "sleep 5m"  
  }, {  
    ...  
  }  
]
```

# HELLO WORLD SCENARIO EXAMPLE

```
{  
  "type" : "load"  
}
```

Will use the default configuration values:

- Use S3 API and port 9020
- Use 1 connection to the default single node @ 127.0.0.1
- Use *Write* load type
- No load limits (infinite load job)
- Use 1MB as the size of the data items
- Use a container (bucket) created automatically



# WEIGHTED LOAD EXAMPLE

```
1.     "type" : "load",
2.     "config" : {
3.         "item" : {
4.             "src" : {
5.                 "file" : [
6.                     null,
7.                     "/tmp/precreated-items-list.csv"
8.                 ]
9.             }
10.        },
11.        "load" : {
12.            "type" : [
13.                "write=20%", "read=80%"
14.            ]
15.        }
16.    }
```

- Performs both write and read operations
- 20% of operations are Write ones
- 80% of operations are Read ones

# WHERE TO GO NEXT

1. Refer to the Mongoose wiki for the *configuration layout* and the detailed scripting engine specification
2. Much *more example scenarios* is available in the Mongoose distribution
3. Ask via email *Mongoose.Support@emc.com*

# OTHER V2.0 FEATURES

## 1.Copy Mode

- *Allows to copy files/directories/objects from the source container/directory to the destination one.*
- *Implemented as an extension of the "Write" load type.*

## 2.Load Limit By Total Size

*It was possible to limit a load job by an item count, a time and a rate in the previous versions. There are the new requirement to make it possible to limit by total processed size. For example, a load job should stop after writing 1TB of a data to the storage.*

# THE ROADMAP FOR 2.X

- GUI Enhancements
- Partial Read
- Centera API Support

# Q & A

# Thank you

EMC<sup>2</sup>®